

```

import pandas as pd

df = pd.read_csv('/content/car_data.csv')

print("First 5 rows of the DataFrame:")
print(df.head())

print("\nDataFrame Info:")
df.info()

print("\nDescriptive Statistics:")
print(df.describe())

print("\nMissing values per column:")
print(df.isnull().sum())

```

First 5 rows of the DataFrame:

|   | Car_Name | Year | Selling_Price | Present_Price | Driven_kms | Fuel_Type |
|---|----------|------|---------------|---------------|------------|-----------|
| 0 | ritz     | 2014 | 3.35          | 5.59          | 27000      | Petrol    |
| 1 | sx4      | 2013 | 4.75          | 9.54          | 43000      | Diesel    |
| 2 | ciaz     | 2017 | 7.25          | 9.85          | 6900       | Petrol    |
| 3 | wagon r  | 2011 | 2.85          | 4.15          | 5200       | Petrol    |
| 4 | swift    | 2014 | 4.60          | 6.87          | 42450      | Diesel    |

|   | Selling_type | Transmission | Owner |
|---|--------------|--------------|-------|
| 0 | Dealer       | Manual       | 0     |
| 1 | Dealer       | Manual       | 0     |
| 2 | Dealer       | Manual       | 0     |
| 3 | Dealer       | Manual       | 0     |
| 4 | Dealer       | Manual       | 0     |

DataFrame Info:

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 301 entries, 0 to 300
Data columns (total 9 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   Car_Name    301 non-null    object 
 1   Year        301 non-null    int64  
 2   Selling_Price 301 non-null  float64
 3   Present_Price 301 non-null  float64
 4   Driven_kms  301 non-null    int64  
 5   Fuel_Type   301 non-null    object 
 6   Selling_type 301 non-null    object 
 7   Transmission 301 non-null    object 
 8   Owner       301 non-null    int64  
dtypes: float64(2), int64(3), object(4)
memory usage: 21.3+ KB

```

Descriptive Statistics:

|       | Year        | Selling_Price | Present_Price | Driven_kms    | Owner      |
|-------|-------------|---------------|---------------|---------------|------------|
| count | 301.000000  | 301.000000    | 301.000000    | 301.000000    | 301.000000 |
| mean  | 2013.627907 | 4.661296      | 7.628472      | 36947.205980  | 0.043189   |
| std   | 2.891554    | 5.082812      | 8.642584      | 38886.883882  | 0.247915   |
| min   | 2003.000000 | 0.100000      | 0.320000      | 500.000000    | 0.000000   |
| 25%   | 2012.000000 | 0.900000      | 1.200000      | 15000.000000  | 0.000000   |
| 50%   | 2014.000000 | 3.600000      | 6.400000      | 32000.000000  | 0.000000   |
| 75%   | 2016.000000 | 6.000000      | 9.900000      | 48767.000000  | 0.000000   |
| max   | 2018.000000 | 35.000000     | 92.600000     | 500000.000000 | 3.000000   |

Missing values per column:

|               | Car_Name | Year | Selling_Price | Present_Price | Driven_kms | Fuel_Type | Selling_type | Transmission | Owner |
|---------------|----------|------|---------------|---------------|------------|-----------|--------------|--------------|-------|
| Car_Name      | 0        |      |               |               |            |           |              |              |       |
| Year          |          | 0    |               |               |            |           |              |              |       |
| Selling_Price |          |      | 0             |               |            |           |              |              |       |
| Present_Price |          |      |               | 0             |            |           |              |              |       |
| Driven_kms    |          |      |               |               | 0          |           |              |              |       |
| Fuel_Type     |          |      |               |               |            | 0         |              |              |       |
| Selling_type  |          |      |               |               |            |           | 0            |              |       |
| Transmission  |          |      |               |               |            |           |              | 0            |       |
| Owner         |          |      |               |               |            |           |              |              | 0     |
|               |          |      |               |               |            |           |              |              |       |

dtype: int64

```

print("Missing values per column (re-verification):")
print(df.isnull().sum())

print("\nUnique values and counts for 'Fuel_Type':")
print(df['Fuel_Type'].value_counts())

```

```

print("\nUnique values and counts for 'Selling_type':")
print(df['Selling_type'].value_counts())

print("\nUnique values and counts for 'Transmission':")
print(df['Transmission'].value_counts())

print("\nDescriptive Statistics for numerical columns (re-verification):")
print(df[['Driven_kms', 'Year', 'Selling_Price', 'Present_Price', 'Owner']].describe())

Missing values per column (re-verification):
Car_Name      0
Year          0
Selling_Price 0
Present_Price 0
Driven_kms    0
Fuel_Type     0
Selling_type   0
Transmission   0
Owner          0
dtype: int64

Unique values and counts for 'Fuel_Type':
Fuel_Type
Petrol    239
Diesel    60
CNG      2
Name: count, dtype: int64

Unique values and counts for 'Selling_type':
Selling_type
Dealer    195
Individual 106
Name: count, dtype: int64

Unique values and counts for 'Transmission':
Transmission
Manual    261
Automatic 40
Name: count, dtype: int64

Descriptive Statistics for numerical columns (re-verification):
   Driven_kms      Year  Selling_Price  Present_Price      Owner
count  301.000000  301.000000  301.000000  301.000000  301.000000
mean   36947.205980 2013.627907   4.661296    7.628472   0.043189
std    38886.883882  2.891554   5.082812   8.642584   0.247915
min    500.000000  2003.000000   0.100000   0.320000   0.000000
25%   15000.000000 2012.000000   0.900000   1.200000   0.000000
50%   32000.000000 2014.000000   3.600000   6.400000   0.000000
75%   48767.000000 2016.000000   6.000000   9.900000   0.000000
max   500000.000000 2018.000000  35.000000  92.600000   3.000000

```

```
df = pd.get_dummies(df, columns=['Fuel_Type', 'Selling_type', 'Transmission'], drop_first=True)
```

```
df.drop('Car_Name', axis=1, inplace=True)
```

```

print("DataFrame after one-hot encoding and dropping 'Car_Name':")
print(df.head())
print("\nDataFrame Info after preprocessing:")
df.info()

```

```

DataFrame after one-hot encoding and dropping 'Car_Name':
   Year  Selling_Price  Present_Price  Driven_kms  Owner  Fuel_Type_Diesel \
0  2014        3.35       5.59     27000      0      False
1  2013        4.75       9.54     43000      0      True
2  2017        7.25       9.85     6900       0     False
3  2011        2.85       4.15     5200       0     False
4  2014        4.60       6.87     42450      0      True

   Fuel_Type_Petrol  Selling_type_Individual  Transmission_Manual
0            True                False                  True
1           False                False                  True
2            True                False                  True
3            True                False                  True
4           False                False                  True

```

```

DataFrame Info after preprocessing:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 301 entries, 0 to 300
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype 

```

```
--  -----
0 Year 301 non-null int64
1 Selling_Price 301 non-null float64
2 Present_Price 301 non-null float64
3 Driven_kms 301 non-null int64
4 Owner 301 non-null int64
5 Fuel_Type_Diesel 301 non-null bool
6 Fuel_Type_Petrol 301 non-null bool
7 Selling_type_Individual 301 non-null bool
8 Transmission_Manual 301 non-null bool
dtypes: bool(4), float64(2), int64(3)
memory usage: 13.1 KB
```

```
df['Age_of_Car'] = 2024 - df['Year']

print("DataFrame with new 'Age_of_Car' feature:")
print(df.head())
```

```
DataFrame with new 'Age_of_Car' feature:
   Year Selling_Price Present_Price Driven_kms Owner Fuel_Type_Diesel \
0  2014      3.35        5.59     27000      0    False
1  2013      4.75        9.54     43000      0     True
2  2017      7.25        9.85     6900       0    False
3  2011      2.85        4.15     5200       0    False
4  2014      4.60        6.87    42450       0     True

   Fuel_Type_Petrol Selling_type_Individual Transmission_Manual Age_of_Car
0            True                False             True          10
1           False                False             True          11
2            True                False             True           7
3            True                False             True          13
4           False                False             True          10
```

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression

# 1. Separate features (X) from the target variable (y)
X = df.drop(['Selling_Price', 'Year'], axis=1)
y = df['Selling_Price']

print("Features (X) shape:", X.shape)
print("Target (y) shape:", y.shape)

# 2. Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

print("\nTraining set shape (X_train, y_train):", X_train.shape, y_train.shape)
print("Testing set shape (X_test, y_test):", X_test.shape, y_test.shape)

# 3. Import LinearRegression (already imported above)

# 4. Instantiate a LinearRegression model
model = LinearRegression()

# 5. Train the model using the training data
model.fit(X_train, y_train)

print("\nLinear Regression model trained successfully.")
```

```
Features (X) shape: (301, 8)
Target (y) shape: (301,)

Training set shape (X_train, y_train): (240, 8) (240,)
Testing set shape (X_test, y_test): (61, 8) (61,)

Linear Regression model trained successfully.
```

```
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error

# Make predictions on the test set
y_pred = model.predict(X_test)

# Evaluate the model
r2 = r2_score(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = mse**0.5 # Calculate RMSE
```

```
print("\nModel Evaluation:")
print(f"R-squared (R2): {r2:.4f}")
print(f"Mean Absolute Error (MAE): {mae:.4f}")
print(f"Mean Squared Error (MSE): {mse:.4f}")
print(f"Root Mean Squared Error (RMSE): {rmse:.4f}")
```

```
Model Evaluation:
R-squared (R2): 0.8489
Mean Absolute Error (MAE): 1.2164
Mean Squared Error (MSE): 3.4813
Root Mean Squared Error (RMSE): 1.8658
```

```
import matplotlib.pyplot as plt

plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred, alpha=0.7)
plt.xlabel('Actual Selling Price (Lakhs)')
plt.ylabel('Predicted Selling Price (Lakhs)')
plt.title('Actual vs. Predicted Selling Prices')
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red', linestyle='--') # Add a diagonal line for perfect
plt.legend(['Predicted Values', 'Perfect Prediction'])
plt.grid(True)
plt.show()
```



